

I decided to treat this as a classification problem by creating a new binary variable `affair` (did the woman have at least one affair?) and trying to predict the classification for each woman.

Dataset

The dataset I chose is the `affairs` dataset that comes with `Statsmodels`. It was derived from a survey of women in 1974 by *Redbook* magazine, in which married women were asked about their participation in extramarital affairs. More information about the study is available in a 1978 paper from the *Journal of Political Economy*.

Description of Variables

The dataset contains 6366 observations of 9 variables:

`rate_marriage`: woman's rating of her marriage (1 = very poor, 5 = very good)

`age`: woman's age

`yrs_married`: number of years married

`children`: number of children

`religious`: woman's rating of how religious she is (1 = not religious, 4 = strongly religious)

`educ`: level of education (9 = grade school, 12 = high school, 14 = some college, 16 = college graduate, 17 = some graduate school, 20 = advanced degree)

occupation: woman's occupation (1 = student, 2 = farming/semi-skilled/unskilled, 3 = "white collar", 4 = teacher/nurse/writer/technician/skilled, 5 = managerial/business, 6 = professional with advanced degree)

occupation_husb: husband's occupation (same coding as above)

affairs: time spent in extra-marital affairs

Code to loading data and modules:

```
import numpy as np

import pandas as pd

import statsmodels.api as sm

import matplotlib.pyplot as plt

from patsy import dmatrices

from sklearn.linear_model

import      LogisticRegression      from

sklearn.cross_validation

import train_test_split from sklearn

import      metrics      from

sklearn.cross_validation

import      cross_val_score      dta      =

sm.datasets.fair.load_pandas().data
```

```

#add "affair" column: 1 represents having affairs, 0
represents not dta['affair'] = (dta.affairs >
0).astype(int)

y, X = dmatrices('affair ~ rate_marriage + age +
    yrs_married + children + \ religious + educ +
    C(occupation) + C(occupation_husb)',
    dta, return_type="dataframe")

X = X.rename(columns =
{'C(occupation)[T.2.0]':'occ_2',
    'C(occupation)[T.3.0]':'occ_3',
    'C(occupation)[T.4.0]':'occ_4',
    'C(occupation)[T.5.0]':'occ_5',
    'C(occupation)[T.6.0]':'occ_6',
    'C(occupation_husb)[T.2.0]':'occ_husb_2',
    'C(occupation_husb)[T.3.0]':'occ_husb_3',
    'C(occupation_husb)[T.4.0]':'occ_husb_4',
    'C(occupation_husb)[T.5.0]':'occ_husb_5',
    'C(occupation_husb)[T.6.0]':'occ_husb_6'})

y = np.ravel(y)

```